

**M2-BigData : GPGPU**  
Chapter 3 – Exercice 2

## Objective

The purpose of this exercise is to convert an RGB image into a gray scale image. The input is an RGB triple of float values and the program will convert that triple to a single float grayscale intensity value. A pseudo-code version of the algorithm is shown below :

```
for i from 0 to height do
  for j from 0 to width do
    idx = i * width + j
    # here channels is 3
    r = input[3*idx]    g = input[3*idx + 1]    b = input[3*idx + 2]
    grayImage[idx] = (0.21*r + 0.71*g + 0.07*b)
  end
end
```

## Instructions

Edit the given code `2-imgTograyscale.cu` to perform the following steps :

- allocate device memory
- copy host memory to device
- initialize thread block and kernel grid dimensions
- invoke CUDA kernel
- copy results from device to host
- deallocate device memory

The execution is performed giving the 2 filenames as program arguments :

```
./2-imgTograyscale inputImg.png outputImg.png
```

The input image must be a colored image. The reading and writing is performed by two functions `read_image_asfloat` and `write_image_fromfloat`. In this exercise, images are represented as RGB or gray with float values.

The CUDA thread grid must be computed from the image dimensions. A common approach is to define squared blocks (called tile), and the grid size accordingly. Mind the arbitrary image dimensions which are generally not a multiple of tile dimensions. Use the given images `Lena.png` then the `Lena1080.png` or any RGB image you want.

## Questions

1. How many floating points operations are being performed in your color conversion kernel? EXPLAIN.
2. How many global memory reads are being performed by your kernel? EXPLAIN.
3. How many global memory writes are being performed by your kernel? EXPLAIN.
4. How do you choose the CUDA thread grid dimensions? EXPLAIN.